

Keras: MLP architectures on MNIST dataset

```
In [0]: # If your keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

```
In [0]: %matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, color=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    print('\n')
    plt.title("Epoch Vs Categorical Cross Entropy Loss")
    fig.canvas.draw()
```

```
In [0]: # The data shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [56]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)" % (X_train.shape[1], X_train.shape[2]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d, %d)" % (X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of test examples : 10000 and each image is of shape (28, 28)

```
In [0]: # If you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of
# 1 * 784
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.sh
ape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2
])
```

```
In [58]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each imag
e is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image
is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [59]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0		
0	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	25	
5	247	127	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	15	
4	170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0		
0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	8	
2	82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	25
3	253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	15	
4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	24
1	225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0		
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	18	
7																			

```

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 249 253 249 64 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 25
3 253 207 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 39 148 229 253 253 253 250 182 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 25
3 253 201 78 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 23 66 213 253 253 253 253 198 81 2 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 18 171 219 253 253 253 253 19
5 80 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 1
6 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]

```

```

In [0]: # Data Normalization
        # if we observe the above matrix each cell is having a value between 0-
        255

```

```
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [61]: # example data point after normlization
print(X_train[0])
```

```
[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.1176471 0.07058824 0.07058824 0.07058824
0.49411765 0.53333333 0.68627451 0.10196078 0.65098039 1.
0.96862745 0.49803922 0.      0.      0.      0.
0.      0.      0.      0.      0.      0.]
```

0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686

0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]

```
In [62]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

Class label of first image : 5
```


After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

1 (a) MLP + ReLU + ADAM with 2 hidden layers

```
In [0]: # https://keras.io/activations/
# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:
from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
In [0]: # Some model parameters
output_dim = 10
input_dim = X_train.shape[1]
batch_size = 128
nb_epoch = 20
```

```
In [65]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_53 (Dense)	(None, 368)	288880
dense_54 (Dense)	(None, 112)	41328
dense_55 (Dense)	(None, 10)	1130

Total params: 331,338
 Trainable params: 331,338
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2490 - acc: 0.9274 - val_loss: 0.1333 - val_acc: 0.9576

Epoch 2/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0956 - acc: 0.9712 - val_loss: 0.0874 - val_acc: 0.9723

Epoch 3/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0618 - acc: 0.9815 - val_loss: 0.0769 - val_acc: 0.9763

Epoch 4/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0417 - acc: 0.9872 - val_loss: 0.0746 - val_acc: 0.9783

Epoch 5/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0308 - acc: 0.9904 - val_loss: 0.0828 - val_acc: 0.9756

Epoch 6/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0227 - acc: 0.9933 - val_loss: 0.0849 - val_acc: 0.9762

Epoch 7/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0189 - acc: 0.9941 - val_loss: 0.0768 - val_acc: 0.9777

Epoch 8/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0186 - acc: 0.9935 - val_loss: 0.0769 - val_acc: 0.9800

```

Epoch 9/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
130 - acc: 0.9960 - val_loss: 0.0709 - val_acc: 0.9818
Epoch 10/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0
091 - acc: 0.9970 - val_loss: 0.0801 - val_acc: 0.9788
Epoch 11/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0
115 - acc: 0.9962 - val_loss: 0.0826 - val_acc: 0.9786
Epoch 12/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
116 - acc: 0.9958 - val_loss: 0.0847 - val_acc: 0.9786
Epoch 13/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
107 - acc: 0.9963 - val_loss: 0.0936 - val_acc: 0.9795
Epoch 14/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
086 - acc: 0.9973 - val_loss: 0.0822 - val_acc: 0.9815
Epoch 15/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
070 - acc: 0.9978 - val_loss: 0.1083 - val_acc: 0.9773
Epoch 16/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
077 - acc: 0.9977 - val_loss: 0.0766 - val_acc: 0.9831
Epoch 17/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
096 - acc: 0.9967 - val_loss: 0.0942 - val_acc: 0.9796
Epoch 18/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
092 - acc: 0.9970 - val_loss: 0.1026 - val_acc: 0.9811
Epoch 19/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
057 - acc: 0.9982 - val_loss: 0.0935 - val_acc: 0.9818
Epoch 20/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
059 - acc: 0.9981 - val_loss: 0.0883 - val_acc: 0.9815

```

```

In [66]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])

```

```

print ('Test accuracy:', score[1])


accuracy_la= score[1]

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08830814581377532
Test accuracy: 0.9815



```

In [67]: w_after = model_relu.get_weights()

         h1_w = w_after[0].flatten().reshape(-1,1)

```

```

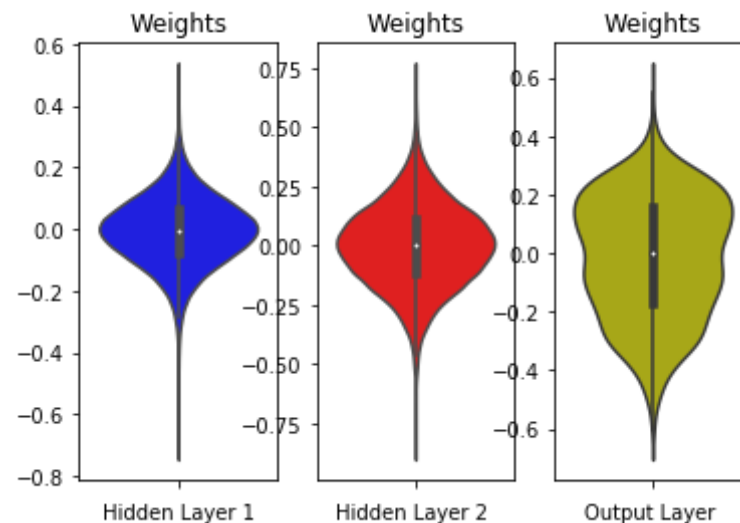
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



1 (b) MLP + ReLU + ADAM + Batch Normalization with 2 hidden layers

```
In [68]: from keras.layers.normalization import BatchNormalization

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_56 (Dense)	(None, 368)	288880
batch_normalization_30 (Batch Normalization)	(None, 368)	1472
dense_57 (Dense)	(None, 112)	41328

```

batch_normalization_31 (Batch Normalization) 448
dense_58 (Dense) (None, 10) 1130
=====
Total params: 333,258
Trainable params: 332,298
Non-trainable params: 960
=====
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 90us/step - loss: 0.2313 - acc: 0.9323 - val_loss: 0.1113 - val_acc: 0.9662
Epoch 2/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0824 - acc: 0.9761 - val_loss: 0.0892 - val_acc: 0.9719
Epoch 3/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0517 - acc: 0.9841 - val_loss: 0.0762 - val_acc: 0.9753
Epoch 4/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0344 - acc: 0.9895 - val_loss: 0.0785 - val_acc: 0.9750
Epoch 5/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0245 - acc: 0.9926 - val_loss: 0.0719 - val_acc: 0.9776
Epoch 6/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0207 - acc: 0.9935 - val_loss: 0.0723 - val_acc: 0.9775
Epoch 7/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0175 - acc: 0.9946 - val_loss: 0.0793 - val_acc: 0.9768
Epoch 8/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0162 - acc: 0.9950 - val_loss: 0.0791 - val_acc: 0.9777
Epoch 9/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0122 - acc: 0.9964 - val_loss: 0.0794 - val_acc: 0.9780
Epoch 10/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0

```

```

132 - acc: 0.9957 - val_loss: 0.0830 - val_acc: 0.9764
Epoch 11/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
103 - acc: 0.9969 - val_loss: 0.0895 - val_acc: 0.9766
Epoch 12/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
114 - acc: 0.9961 - val_loss: 0.0838 - val_acc: 0.9775
Epoch 13/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
104 - acc: 0.9967 - val_loss: 0.0846 - val_acc: 0.9769
Epoch 14/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
110 - acc: 0.9962 - val_loss: 0.0744 - val_acc: 0.9811
Epoch 15/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
068 - acc: 0.9979 - val_loss: 0.0704 - val_acc: 0.9811
Epoch 16/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
066 - acc: 0.9979 - val_loss: 0.0875 - val_acc: 0.9782
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
085 - acc: 0.9972 - val_loss: 0.0985 - val_acc: 0.9759
Epoch 18/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
089 - acc: 0.9969 - val_loss: 0.0764 - val_acc: 0.9807
Epoch 19/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
066 - acc: 0.9977 - val_loss: 0.0727 - val_acc: 0.9823
Epoch 20/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
050 - acc: 0.9986 - val_loss: 0.0749 - val_acc: 0.9834

```

```

In [69]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_lb= score[1]

fig, ax = plt.subplots(1,1)

```



```


ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07494691710416282

Test accuracy: 0.9834



```

In [70]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()

```

```
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



1 (c) MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 2 hidden layers

```
In [71]: from keras.layers import Dropout
```

```

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_59 (Dense)	(None, 368)	288880
batch_normalization_32 (Batch Normalization)	(None, 368)	1472
dropout_20 (Dropout)	(None, 368)	0
dense_60 (Dense)	(None, 112)	41328
batch_normalization_33 (Batch Normalization)	(None, 112)	448

dropout_21 (Dropout)	(None, 112)	0
dense_61 (Dense)	(None, 10)	1130

Total params: 333,258
 Trainable params: 332,298
 Non-trainable params: 960

None
 Train on 60000 samples, validate on 10000 samples
 Epoch 1/20
 60000/60000 [=====] - 6s 96us/step - loss: 0.4490 - acc: 0.8626 - val_loss: 0.1546 - val_acc: 0.9530
 Epoch 2/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.2083 - acc: 0.9376 - val_loss: 0.1117 - val_acc: 0.9642
 Epoch 3/20
 60000/60000 [=====] - 3s 53us/step - loss: 0.1589 - acc: 0.9523 - val_loss: 0.0880 - val_acc: 0.9723
 Epoch 4/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.1358 - acc: 0.9577 - val_loss: 0.0810 - val_acc: 0.9744
 Epoch 5/20
 60000/60000 [=====] - 3s 53us/step - loss: 0.1159 - acc: 0.9641 - val_loss: 0.0771 - val_acc: 0.9752
 Epoch 6/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.1046 - acc: 0.9676 - val_loss: 0.0730 - val_acc: 0.9771
 Epoch 7/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.0933 - acc: 0.9715 - val_loss: 0.0680 - val_acc: 0.9796
 Epoch 8/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.0842 - acc: 0.9734 - val_loss: 0.0681 - val_acc: 0.9789
 Epoch 9/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.0811 - acc: 0.9747 - val_loss: 0.0616 - val_acc: 0.9814
 Epoch 10/20

```

60000/60000 [=====] - 3s 53us/step - loss: 0.0
729 - acc: 0.9765 - val_loss: 0.0610 - val_acc: 0.9815
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
676 - acc: 0.9780 - val_loss: 0.0583 - val_acc: 0.9824
Epoch 12/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
634 - acc: 0.9802 - val_loss: 0.0599 - val_acc: 0.9818
Epoch 13/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
613 - acc: 0.9799 - val_loss: 0.0571 - val_acc: 0.9825
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
587 - acc: 0.9816 - val_loss: 0.0591 - val_acc: 0.9817
Epoch 15/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
536 - acc: 0.9825 - val_loss: 0.0560 - val_acc: 0.9825
Epoch 16/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
512 - acc: 0.9835 - val_loss: 0.0568 - val_acc: 0.9829
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
507 - acc: 0.9835 - val_loss: 0.0579 - val_acc: 0.9823
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
517 - acc: 0.9832 - val_loss: 0.0592 - val_acc: 0.9826
Epoch 19/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
453 - acc: 0.9852 - val_loss: 0.0599 - val_acc: 0.9832
Epoch 20/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
460 - acc: 0.9853 - val_loss: 0.0610 - val_acc: 0.9827

```

```

In [72]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_1c= score[1]


```

```
fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.061045588629494885

Test accuracy: 0.9827



```
In [73]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



1 (d) MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 2 hidden layers

```
In [74]: from keras.layers import Dropout

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_62 (Dense)	(None, 368)	288880
batch_normalization_34 (Batch Normalization)	(None, 368)	1472
dropout_22 (Dropout)	(None, 368)	0
dense_63 (Dense)	(None, 112)	41328

batch_normalization_35 (Batch Normalization)	(None, 112)	448
dropout_23 (Dropout)	(None, 112)	0
dense_64 (Dense)	(None, 10)	1130

Total params: 333,258
 Trainable params: 332,298
 Non-trainable params: 960

None
 Train on 60000 samples, validate on 10000 samples
 Epoch 1/20
 60000/60000 [=====] - 6s 101us/step - loss: 0.5493 - acc: 0.8347 - val_loss: 0.1784 - val_acc: 0.9437
 Epoch 2/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.2607 - acc: 0.9229 - val_loss: 0.1306 - val_acc: 0.9593
 Epoch 3/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.2022 - acc: 0.9396 - val_loss: 0.1035 - val_acc: 0.9676
 Epoch 4/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.1690 - acc: 0.9494 - val_loss: 0.0926 - val_acc: 0.9707
 Epoch 5/20
 60000/60000 [=====] - 3s 54us/step - loss: 0.1471 - acc: 0.9562 - val_loss: 0.0864 - val_acc: 0.9722
 Epoch 6/20
 60000/60000 [=====] - 3s 54us/step - loss: 0.1338 - acc: 0.9598 - val_loss: 0.0793 - val_acc: 0.9753
 Epoch 7/20
 60000/60000 [=====] - 3s 54us/step - loss: 0.1219 - acc: 0.9634 - val_loss: 0.0767 - val_acc: 0.9763
 Epoch 8/20
 60000/60000 [=====] - 3s 55us/step - loss: 0.1128 - acc: 0.9658 - val_loss: 0.0690 - val_acc: 0.9785
 Epoch 9/20
 60000/60000 [=====] - 3s 54us/step - loss: 0.1105 - acc: 0.9663 - val_loss: 0.0746 - val_acc: 0.9779

```

Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
998 - acc: 0.9686 - val_loss: 0.0699 - val_acc: 0.9782
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
937 - acc: 0.9707 - val_loss: 0.0681 - val_acc: 0.9794
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
897 - acc: 0.9726 - val_loss: 0.0666 - val_acc: 0.9795
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
838 - acc: 0.9749 - val_loss: 0.0625 - val_acc: 0.9806
Epoch 14/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
836 - acc: 0.9744 - val_loss: 0.0626 - val_acc: 0.9830
Epoch 15/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
746 - acc: 0.9762 - val_loss: 0.0609 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
743 - acc: 0.9766 - val_loss: 0.0595 - val_acc: 0.9818
Epoch 17/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
698 - acc: 0.9780 - val_loss: 0.0625 - val_acc: 0.9822
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
693 - acc: 0.9789 - val_loss: 0.0558 - val_acc: 0.9844
Epoch 19/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0
653 - acc: 0.9790 - val_loss: 0.0586 - val_acc: 0.9834
Epoch 20/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0
648 - acc: 0.9797 - val_loss: 0.0586 - val_acc: 0.9834

```

```

In [75]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])


accuracy_ld= score[1]

```

```
fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0585743734298856
 Test accuacy: 0.9834



```
In [76]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



1 (e) MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 2 hidden layers

```
In [77]: from keras.layers import Dropout

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====		
dense_65 (Dense)	(None, 368)	288880
batch_normalization_36 (Batch Normalization)	(None, 368)	1472
dropout_24 (Dropout)	(None, 368)	0
dense_66 (Dense)	(None, 112)	41328

batch_normalization_37 (Batch Normalization)	(None, 112)	448
dropout_25 (Dropout)	(None, 112)	0
dense_67 (Dense)	(None, 10)	1130

Total params: 333,258
 Trainable params: 332,298
 Non-trainable params: 960

None
 Train on 60000 samples, validate on 10000 samples
 Epoch 1/20
 60000/60000 [=====] - 6s 106us/step - loss: 0.7216 - acc: 0.7808 - val_loss: 0.2077 - val_acc: 0.9364
 Epoch 2/20
 60000/60000 [=====] - 3s 54us/step - loss: 0.3329 - acc: 0.9003 - val_loss: 0.1592 - val_acc: 0.9515
 Epoch 3/20
 60000/60000 [=====] - 3s 55us/step - loss: 0.2606 - acc: 0.9235 - val_loss: 0.1266 - val_acc: 0.9605
 Epoch 4/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.2161 - acc: 0.9358 - val_loss: 0.1145 - val_acc: 0.9644
 Epoch 5/20
 60000/60000 [=====] - 3s 54us/step - loss: 0.1976 - acc: 0.9424 - val_loss: 0.1005 - val_acc: 0.9698
 Epoch 6/20
 60000/60000 [=====] - 3s 53us/step - loss: 0.1759 - acc: 0.9487 - val_loss: 0.0933 - val_acc: 0.9716
 Epoch 7/20
 60000/60000 [=====] - 3s 53us/step - loss: 0.1645 - acc: 0.9523 - val_loss: 0.0892 - val_acc: 0.9733
 Epoch 8/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.1506 - acc: 0.9564 - val_loss: 0.0877 - val_acc: 0.9742
 Epoch 9/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.1

```

396 - acc: 0.9588 - val_loss: 0.0804 - val_acc: 0.9762
Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1
373 - acc: 0.9600 - val_loss: 0.0793 - val_acc: 0.9756
Epoch 11/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1
262 - acc: 0.9622 - val_loss: 0.0760 - val_acc: 0.9776
Epoch 12/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1
181 - acc: 0.9655 - val_loss: 0.0748 - val_acc: 0.9782
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1
187 - acc: 0.9652 - val_loss: 0.0731 - val_acc: 0.9783
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1
100 - acc: 0.9676 - val_loss: 0.0716 - val_acc: 0.9784
Epoch 15/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1
112 - acc: 0.9665 - val_loss: 0.0698 - val_acc: 0.9786
Epoch 16/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1
014 - acc: 0.9699 - val_loss: 0.0694 - val_acc: 0.9795
Epoch 17/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0
997 - acc: 0.9696 - val_loss: 0.0697 - val_acc: 0.9793
Epoch 18/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
972 - acc: 0.9703 - val_loss: 0.0702 - val_acc: 0.9791
Epoch 19/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0
965 - acc: 0.9713 - val_loss: 0.0656 - val_acc: 0.9802
Epoch 20/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
927 - acc: 0.9719 - val_loss: 0.0630 - val_acc: 0.9811

```

```

In [78]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

```

```


accuracy_1e= score[1]

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0630017356009921
Test accuacy: 0.9811



```

In [79]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```




```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



2 (a) MLP + ReLU + ADAM with 3 hidden layers

```
In [80]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_68 (Dense)	(None, 368)	288880
dense_69 (Dense)	(None, 112)	41328
dense_70 (Dense)	(None, 54)	6102
dense_71 (Dense)	(None, 10)	550
=====	=====	=====
Total params: 336,860		
Trainable params: 336,860		
Non-trainable params: 0		
None		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 84us/step - loss: 0.2
749 - acc: 0.9166 - val_loss: 0.1294 - val_acc: 0.9611
Epoch 2/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
950 - acc: 0.9714 - val_loss: 0.0870 - val_acc: 0.9730
Epoch 3/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
609 - acc: 0.9812 - val_loss: 0.0900 - val_acc: 0.9723
Epoch 4/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
438 - acc: 0.9860 - val_loss: 0.0870 - val_acc: 0.9714
Epoch 5/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
333 - acc: 0.9888 - val_loss: 0.0900 - val_acc: 0.9724
Epoch 6/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
259 - acc: 0.9920 - val_loss: 0.0782 - val_acc: 0.9774
Epoch 7/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
225 - acc: 0.9927 - val_loss: 0.0987 - val_acc: 0.9742
Epoch 8/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
204 - acc: 0.9933 - val_loss: 0.0864 - val_acc: 0.9771
Epoch 9/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
185 - acc: 0.9937 - val_loss: 0.0984 - val_acc: 0.9744
Epoch 10/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
146 - acc: 0.9952 - val_loss: 0.0883 - val_acc: 0.9782
Epoch 11/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0
140 - acc: 0.9955 - val_loss: 0.0915 - val_acc: 0.9778
Epoch 12/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0
154 - acc: 0.9948 - val_loss: 0.0852 - val_acc: 0.9805
```

```

Epoch 13/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0
117 - acc: 0.9960 - val_loss: 0.0974 - val_acc: 0.9799
Epoch 14/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
090 - acc: 0.9969 - val_loss: 0.0894 - val_acc: 0.9804
Epoch 15/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
154 - acc: 0.9953 - val_loss: 0.1019 - val_acc: 0.9776
Epoch 16/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
142 - acc: 0.9955 - val_loss: 0.0876 - val_acc: 0.9811
Epoch 17/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
090 - acc: 0.9972 - val_loss: 0.0860 - val_acc: 0.9807
Epoch 18/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
051 - acc: 0.9984 - val_loss: 0.1003 - val_acc: 0.9794
Epoch 19/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
118 - acc: 0.9962 - val_loss: 0.0873 - val_acc: 0.9801
Epoch 20/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
089 - acc: 0.9971 - val_loss: 0.0966 - val_acc: 0.9791

```

```

In [81]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_2a= score[1]

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')


#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']

```

```
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09661503153744422

Test accuracy: 0.9791




```
In [82]: w_after = model_relu.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
h3_w = w_after[4].flatten().reshape(-1,1)  
out_w = w_after[6].flatten().reshape(-1,1)  
  
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 4, 1)  
plt.title("Weights")  
ax = sns.violinplot(y=h1_w,color='b')  
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2 (b) MLP + ReLU + ADAM + Batch Normalization with 3 hidden layers

```
In [83]: from keras.layers.normalization import BatchNormalization
```

```

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_72 (Dense)	(None, 368)	288880
batch_normalization_38 (Batch Normalization)	(None, 368)	1472
dense_73 (Dense)	(None, 112)	41328

batch_normalization_39 (Batch Normalization)	(None, 112)	448
dense_74 (Dense)	(None, 54)	6102
batch_normalization_40 (Batch Normalization)	(None, 54)	216
dense_75 (Dense)	(None, 10)	550

=====

Total params: 338,996
Trainable params: 337,928
Non-trainable params: 1,068

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 7s 118us/step - loss: 0.2583 - acc: 0.9254 - val_loss: 0.1199 - val_acc: 0.9635

Epoch 2/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0862 - acc: 0.9742 - val_loss: 0.0901 - val_acc: 0.9711

Epoch 3/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0546 - acc: 0.9839 - val_loss: 0.0881 - val_acc: 0.9717

Epoch 4/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0387 - acc: 0.9878 - val_loss: 0.0782 - val_acc: 0.9754

Epoch 5/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0290 - acc: 0.9911 - val_loss: 0.0988 - val_acc: 0.9701

Epoch 6/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0244 - acc: 0.9921 - val_loss: 0.0887 - val_acc: 0.9748

Epoch 7/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0197 - acc: 0.9936 - val_loss: 0.0777 - val_acc: 0.9782

Epoch 8/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0


```
164 - acc: 0.9949 - val_loss: 0.0902 - val_acc: 0.9753
Epoch 9/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0
161 - acc: 0.9944 - val_loss: 0.0962 - val_acc: 0.9740
Epoch 10/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0
153 - acc: 0.9950 - val_loss: 0.0831 - val_acc: 0.9787
Epoch 11/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0
140 - acc: 0.9954 - val_loss: 0.0872 - val_acc: 0.9767
Epoch 12/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0
141 - acc: 0.9950 - val_loss: 0.0850 - val_acc: 0.9774
Epoch 13/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0
137 - acc: 0.9953 - val_loss: 0.0790 - val_acc: 0.9803
Epoch 14/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0
120 - acc: 0.9959 - val_loss: 0.0760 - val_acc: 0.9802
Epoch 15/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0
100 - acc: 0.9967 - val_loss: 0.0879 - val_acc: 0.9787
Epoch 16/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0
084 - acc: 0.9973 - val_loss: 0.0833 - val_acc: 0.9798
Epoch 17/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0
060 - acc: 0.9982 - val_loss: 0.0757 - val_acc: 0.9805
Epoch 18/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
092 - acc: 0.9969 - val_loss: 0.0917 - val_acc: 0.9777
Epoch 19/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0
094 - acc: 0.9965 - val_loss: 0.0776 - val_acc: 0.9810
Epoch 20/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0
096 - acc: 0.9966 - val_loss: 0.0940 - val_acc: 0.9790
```

```
In [84]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
```

```

print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_2b= score[1]


fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09398488705815872

Test accuracy: 0.979



In [85]: w_after = model_relu.get_weights()


```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2 (c) MLP + ReLU + ADAM + Batch Normalization + Dropout (0.4) with 3 hidden layers

```
In [86]: from keras.layers.normalization import BatchNormalization

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
```

```

model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_76 (Dense)	(None, 368)	288880
batch_normalization_41 (Batch Normalization)	(None, 368)	1472
dropout_26 (Dropout)	(None, 368)	0
dense_77 (Dense)	(None, 112)	41328
batch_normalization_42 (Batch Normalization)	(None, 112)	448
dropout_27 (Dropout)	(None, 112)	0
dense_78 (Dense)	(None, 54)	6102
batch_normalization_43 (Batch Normalization)	(None, 54)	216
dropout_28 (Dropout)	(None, 54)	0
dense_79 (Dense)	(None, 10)	550

```
=====
Total params: 338,996
Trainable params: 337,928
Non-trainable params: 1,068

None

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 126us/step - loss: 0.
6901 - acc: 0.7904 - val_loss: 0.1926 - val_acc: 0.9411
Epoch 2/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2
870 - acc: 0.9175 - val_loss: 0.1392 - val_acc: 0.9573
Epoch 3/20
60000/60000 [=====] - 4s 68us/step - loss: 0.2
203 - acc: 0.9373 - val_loss: 0.1100 - val_acc: 0.9671
Epoch 4/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1
741 - acc: 0.9496 - val_loss: 0.0996 - val_acc: 0.9696
Epoch 5/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1
518 - acc: 0.9565 - val_loss: 0.0983 - val_acc: 0.9714
Epoch 6/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
383 - acc: 0.9602 - val_loss: 0.0861 - val_acc: 0.9763
Epoch 7/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
245 - acc: 0.9644 - val_loss: 0.0811 - val_acc: 0.9761
Epoch 8/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
138 - acc: 0.9667 - val_loss: 0.0797 - val_acc: 0.9776
Epoch 9/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
086 - acc: 0.9688 - val_loss: 0.0762 - val_acc: 0.9780
Epoch 10/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0
984 - acc: 0.9719 - val_loss: 0.0736 - val_acc: 0.9769
```

```

Epoch 11/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0
937 - acc: 0.9729 - val_loss: 0.0692 - val_acc: 0.9786
Epoch 12/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0
904 - acc: 0.9738 - val_loss: 0.0652 - val_acc: 0.9806
Epoch 13/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0
847 - acc: 0.9752 - val_loss: 0.0634 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0
820 - acc: 0.9760 - val_loss: 0.0675 - val_acc: 0.9805
Epoch 15/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0
751 - acc: 0.9776 - val_loss: 0.0703 - val_acc: 0.9789
Epoch 16/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0
712 - acc: 0.9786 - val_loss: 0.0681 - val_acc: 0.9811
Epoch 17/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0
679 - acc: 0.9801 - val_loss: 0.0661 - val_acc: 0.9809
Epoch 18/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0
640 - acc: 0.9810 - val_loss: 0.0618 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0
653 - acc: 0.9805 - val_loss: 0.0644 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0
624 - acc: 0.9816 - val_loss: 0.0634 - val_acc: 0.9821

```

```

In [87]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_2c= score[1]

fig, ax = plt.subplots(1,1)

```

```


ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06336671170063783

Test accuracy: 0.9821



```

In [88]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

```



```


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



2 (d) MLP + ReLU + ADAM + Batch Normalization + Dropout (0.5) with 3 hidden layers

```
In [89]: from keras.layers.normalization import BatchNormalization

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_80 (Dense)	(None, 368)	288880
batch_normalization_44 (Batch Normalization)	(None, 368)	1472
dropout_29 (Dropout)	(None, 368)	0
dense_81 (Dense)	(None, 112)	41328
batch_normalization_45 (Batch Normalization)	(None, 112)	448
dropout_30 (Dropout)	(None, 112)	0
dense_82 (Dense)	(None, 54)	6102
batch_normalization_46 (Batch Normalization)	(None, 54)	216
dropout_31 (Dropout)	(None, 54)	0
dense_83 (Dense)	(None, 10)	550
Total params: 338,996		
Trainable params: 337,928		
Non-trainable params: 1,068		
None		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 141us/step - loss: 0.9417 - acc: 0.7092 - val_loss: 0.2406 - val_acc: 0.9293

Epoch 2/20

60000/60000 [=====] - 4s 68us/step - loss: 0.3996 - acc: 0.8859 - val_loss: 0.1663 - val_acc: 0.9507

Epoch 3/20

60000/60000 [=====] - 4s 72us/step - loss: 0.2950 - acc: 0.9180 - val_loss: 0.1343 - val_acc: 0.9609

```
Epoch 4/20
60000/60000 [=====] - 4s 72us/step - loss: 0.2
490 - acc: 0.9312 - val_loss: 0.1180 - val_acc: 0.9656
Epoch 5/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2
180 - acc: 0.9401 - val_loss: 0.1118 - val_acc: 0.9679
Epoch 6/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
913 - acc: 0.9472 - val_loss: 0.1003 - val_acc: 0.9702
Epoch 7/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
728 - acc: 0.9525 - val_loss: 0.0952 - val_acc: 0.9712
Epoch 8/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
624 - acc: 0.9561 - val_loss: 0.0917 - val_acc: 0.9756
Epoch 9/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
495 - acc: 0.9598 - val_loss: 0.0851 - val_acc: 0.9744
Epoch 10/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1
421 - acc: 0.9620 - val_loss: 0.0855 - val_acc: 0.9751
Epoch 11/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
307 - acc: 0.9646 - val_loss: 0.0793 - val_acc: 0.9775
Epoch 12/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
269 - acc: 0.9655 - val_loss: 0.0770 - val_acc: 0.9783
Epoch 13/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
200 - acc: 0.9677 - val_loss: 0.0755 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
167 - acc: 0.9679 - val_loss: 0.0746 - val_acc: 0.9791
Epoch 15/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
137 - acc: 0.9693 - val_loss: 0.0752 - val_acc: 0.9787
Epoch 16/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
043 - acc: 0.9703 - val_loss: 0.0736 - val_acc: 0.9810
```

```
Epoch 17/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0
999 - acc: 0.9720 - val_loss: 0.0748 - val_acc: 0.9790
Epoch 18/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
000 - acc: 0.9715 - val_loss: 0.0725 - val_acc: 0.9812
Epoch 19/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0
942 - acc: 0.9738 - val_loss: 0.0723 - val_acc: 0.9803
Epoch 20/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0
924 - acc: 0.9747 - val_loss: 0.0742 - val_acc: 0.9804
```


```
In [90]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_2d= score[1]

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.07416196312536485
Test accuracy: 0.9804
```



```
In [91]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

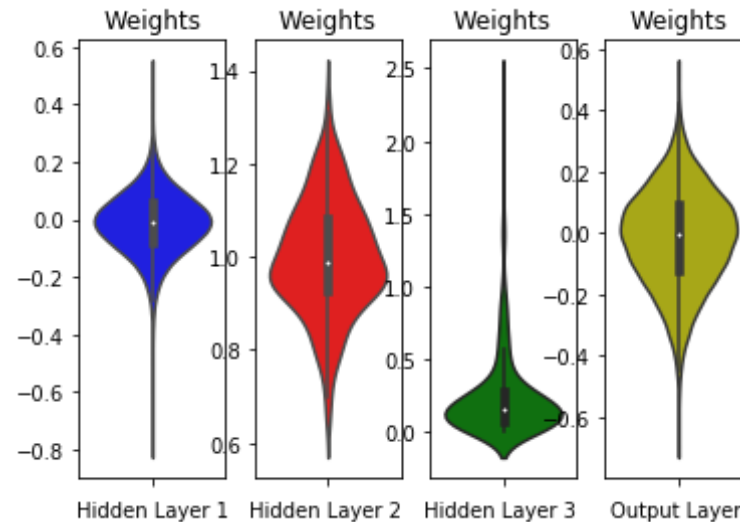
plt.subplot(1, 4, 3)
plt.title("Weights")
```

```

ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



2 (e) MLP + ReLU + ADAM + Batch Normalization + Dropout (0.6) with 3 hidden layers

```

In [92]: from keras.layers.normalization import BatchNormalization

model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.073 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.133 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

```

```

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_84 (Dense)	(None, 368)	288880
batch_normalization_47 (Batch Normalization)	(None, 368)	1472
dropout_32 (Dropout)	(None, 368)	0
dense_85 (Dense)	(None, 112)	41328
batch_normalization_48 (Batch Normalization)	(None, 112)	448
dropout_33 (Dropout)	(None, 112)	0

dense_86 (Dense)	(None, 54)	6102
batch_normalization_49 (Batch Normalization)	(None, 54)	216
dropout_34 (Dropout)	(None, 54)	0
dense_87 (Dense)	(None, 10)	550
=====		
Total params: 338,996		
Trainable params: 337,928		
Non-trainable params: 1,068		
None		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 141us/step - loss: 1.3064 - acc: 0.5828 - val_loss: 0.3388 - val_acc: 0.9123

Epoch 2/20

60000/60000 [=====] - 4s 67us/step - loss: 0.5988 - acc: 0.8253 - val_loss: 0.2308 - val_acc: 0.9316

Epoch 3/20

60000/60000 [=====] - 4s 66us/step - loss: 0.4305 - acc: 0.8799 - val_loss: 0.1790 - val_acc: 0.9484

Epoch 4/20

60000/60000 [=====] - 4s 67us/step - loss: 0.3526 - acc: 0.9046 - val_loss: 0.1603 - val_acc: 0.9532

Epoch 5/20

60000/60000 [=====] - 4s 66us/step - loss: 0.3052 - acc: 0.9188 - val_loss: 0.1424 - val_acc: 0.9589

Epoch 6/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2752 - acc: 0.9273 - val_loss: 0.1270 - val_acc: 0.9635

Epoch 7/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2475 - acc: 0.9349 - val_loss: 0.1184 - val_acc: 0.9662

Epoch 8/20

```
60000/60000 [=====] - 4s 67us/step - loss: 0.2
309 - acc: 0.9393 - val_loss: 0.1120 - val_acc: 0.9685
Epoch 9/20
60000/60000 [=====] - 4s 68us/step - loss: 0.2
143 - acc: 0.9444 - val_loss: 0.1107 - val_acc: 0.9694
Epoch 10/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2
018 - acc: 0.9479 - val_loss: 0.1044 - val_acc: 0.9712
Epoch 11/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
901 - acc: 0.9498 - val_loss: 0.0953 - val_acc: 0.9733
Epoch 12/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
802 - acc: 0.9529 - val_loss: 0.0974 - val_acc: 0.9733
Epoch 13/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
741 - acc: 0.9550 - val_loss: 0.0917 - val_acc: 0.9748
Epoch 14/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
709 - acc: 0.9560 - val_loss: 0.0904 - val_acc: 0.9747
Epoch 15/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
628 - acc: 0.9581 - val_loss: 0.0911 - val_acc: 0.9747
Epoch 16/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
605 - acc: 0.9586 - val_loss: 0.0853 - val_acc: 0.9771
Epoch 17/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
537 - acc: 0.9596 - val_loss: 0.0836 - val_acc: 0.9777
Epoch 18/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
474 - acc: 0.9607 - val_loss: 0.0825 - val_acc: 0.9782
Epoch 19/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1
389 - acc: 0.9633 - val_loss: 0.0853 - val_acc: 0.9789
Epoch 20/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1
364 - acc: 0.9646 - val_loss: 0.0825 - val_acc: 0.9786
```

```

In [93]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

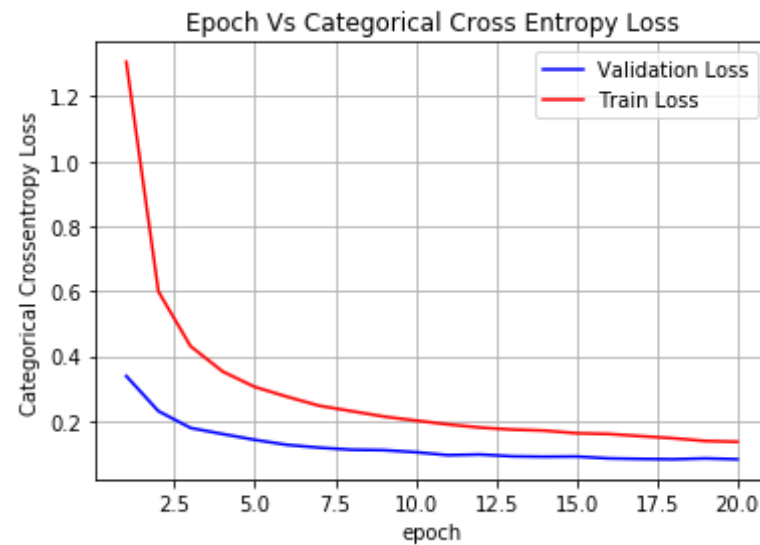
accuracy_2e= score[1]

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08248090120132547
Test accuracy: 0.9786



```

In [94]: w_after = model_relu.get_weights()

```


```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 (a) MLP + ReLU + ADAM with 5 hidden layers

```
In [95]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(Dense(28, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.267, seed=None)))
model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.353, seed=None)))
```

```

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_88 (Dense)	(None, 368)	288880
dense_89 (Dense)	(None, 112)	41328
dense_90 (Dense)	(None, 54)	6102
dense_91 (Dense)	(None, 28)	1540
dense_92 (Dense)	(None, 16)	464
dense_93 (Dense)	(None, 10)	170
Total params: 338,484		
Trainable params: 338,484		
Non-trainable params: 0		
None		

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 110us/step - loss: 0.3656 - acc: 0.8884 - val_loss: 0.1626 - val_acc: 0.9530
Epoch 2/20
60000/60000 [=====] - 3s 42us/step - loss: 0.1205 - acc: 0.9650 - val_loss: 0.1107 - val_acc: 0.9671
Epoch 3/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0

```

```
797 - acc: 0.9756 - val_loss: 0.1117 - val_acc: 0.9679
Epoch 4/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
577 - acc: 0.9824 - val_loss: 0.0969 - val_acc: 0.9733
Epoch 5/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0
470 - acc: 0.9854 - val_loss: 0.0905 - val_acc: 0.9742
Epoch 6/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
384 - acc: 0.9878 - val_loss: 0.0874 - val_acc: 0.9756
Epoch 7/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
307 - acc: 0.9901 - val_loss: 0.1153 - val_acc: 0.9696
Epoch 8/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
272 - acc: 0.9909 - val_loss: 0.0947 - val_acc: 0.9748
Epoch 9/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
268 - acc: 0.9912 - val_loss: 0.1115 - val_acc: 0.9725
Epoch 10/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
213 - acc: 0.9929 - val_loss: 0.1055 - val_acc: 0.9763
Epoch 11/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
216 - acc: 0.9929 - val_loss: 0.0983 - val_acc: 0.9764
Epoch 12/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
157 - acc: 0.9948 - val_loss: 0.1091 - val_acc: 0.9753
Epoch 13/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0
155 - acc: 0.9947 - val_loss: 0.0988 - val_acc: 0.9766
Epoch 14/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0
206 - acc: 0.9936 - val_loss: 0.1010 - val_acc: 0.9774
Epoch 15/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0
137 - acc: 0.9956 - val_loss: 0.0986 - val_acc: 0.9787
Epoch 16/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
```

```
140 - acc: 0.9951 - val_loss: 0.0991 - val_acc: 0.9799
Epoch 17/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
139 - acc: 0.9956 - val_loss: 0.0978 - val_acc: 0.9792
Epoch 18/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0
092 - acc: 0.9970 - val_loss: 0.1148 - val_acc: 0.9773
Epoch 19/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0
126 - acc: 0.9960 - val_loss: 0.1070 - val_acc: 0.9778
Epoch 20/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0
142 - acc: 0.9957 - val_loss: 0.1238 - val_acc: 0.9763
```


```
In [96]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

accuracy_3a= score[1]

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.12384451112879369
Test accuracy: 0.9763
```

```
In [97]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Layer 1')

plt.subplot(1, 6, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 2 ')
```

```

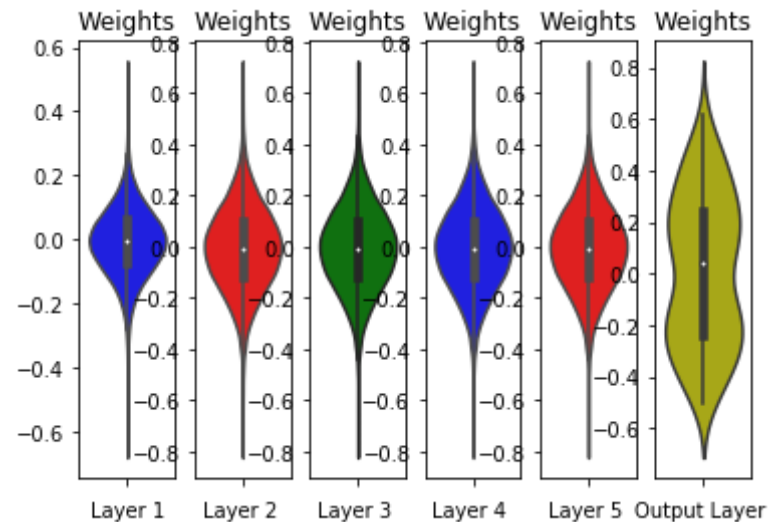
plt.subplot(1, 6, 3)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='g')
plt.xlabel('Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



3 (b) MLP + ReLU + ADAM + Batch Normalization with 5 hidden layers

```
In [98]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(28, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.267, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.353, seed=None)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_94 (Dense)	(None, 368)	288880
batch_normalization_50 (Batch Normalization)	(None, 368)	1472
dense_95 (Dense)	(None, 112)	41328
batch_normalization_51 (Batch Normalization)	(None, 112)	448
dense_96 (Dense)	(None, 54)	6102
batch_normalization_52 (Batch Normalization)	(None, 54)	216
dense_97 (Dense)	(None, 28)	1540
batch_normalization_53 (Batch Normalization)	(None, 28)	112
dense_98 (Dense)	(None, 16)	464
batch_normalization_54 (Batch Normalization)	(None, 16)	64
dense_99 (Dense)	(None, 10)	170
=====	=====	=====
Total params: 340,796		
Trainable params: 339,640		
Non-trainable params: 1,156		
None		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 173us/step - loss: 0.4396 - acc: 0.8921 - val_loss: 0.1558 - val_acc: 0.9576
Epoch 2/20
```

```
60000/60000 [=====] - 5s 89us/step - loss: 0.1
295 - acc: 0.9650 - val_loss: 0.1105 - val_acc: 0.9679
Epoch 3/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
822 - acc: 0.9767 - val_loss: 0.0968 - val_acc: 0.9718
Epoch 4/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
592 - acc: 0.9830 - val_loss: 0.0971 - val_acc: 0.9707
Epoch 5/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0
496 - acc: 0.9848 - val_loss: 0.0857 - val_acc: 0.9755
Epoch 6/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
394 - acc: 0.9879 - val_loss: 0.0959 - val_acc: 0.9732
Epoch 7/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
324 - acc: 0.9905 - val_loss: 0.0806 - val_acc: 0.9759
Epoch 8/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
271 - acc: 0.9916 - val_loss: 0.0825 - val_acc: 0.9756
Epoch 9/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
271 - acc: 0.9912 - val_loss: 0.0854 - val_acc: 0.9780
Epoch 10/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
239 - acc: 0.9919 - val_loss: 0.0843 - val_acc: 0.9777
Epoch 11/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
225 - acc: 0.9926 - val_loss: 0.0984 - val_acc: 0.9734
Epoch 12/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
205 - acc: 0.9934 - val_loss: 0.0934 - val_acc: 0.9764
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
179 - acc: 0.9940 - val_loss: 0.0851 - val_acc: 0.9782
Epoch 14/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
145 - acc: 0.9954 - val_loss: 0.0825 - val_acc: 0.9781
Epoch 15/20
```

```

60000/60000 [=====] - 5s 91us/step - loss: 0.0
173 - acc: 0.9944 - val_loss: 0.0872 - val_acc: 0.9781
Epoch 16/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
148 - acc: 0.9952 - val_loss: 0.0839 - val_acc: 0.9780
Epoch 17/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
118 - acc: 0.9961 - val_loss: 0.0768 - val_acc: 0.9799
Epoch 18/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
138 - acc: 0.9953 - val_loss: 0.0865 - val_acc: 0.9789
Epoch 19/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
137 - acc: 0.9956 - val_loss: 0.0826 - val_acc: 0.9790
Epoch 20/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
110 - acc: 0.9966 - val_loss: 0.0853 - val_acc: 0.9794

```

```

In [99]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])

```

```

accuracy_3b= score[1]

```

```

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

```

```


#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.08531874578401911
Test accuracy: 0.9794

```



```
In [100]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Layer 1')


plt.subplot(1, 6, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 2 ')
```

```
plt.subplot(1, 6, 3)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='g')
plt.xlabel('Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 (c) MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 5 hidden layers

```
In [101]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(28, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.267, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.353, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))

model_relu.add(Dense(output_dim, activation='softmax'))
```

```
print(model_relu.summary())
print('\n')
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_100 (Dense)	(None, 368)	288880
batch_normalization_55 (Batch Normalization)	(None, 368)	1472
dropout_35 (Dropout)	(None, 368)	0
dense_101 (Dense)	(None, 112)	41328
batch_normalization_56 (Batch Normalization)	(None, 112)	448
dropout_36 (Dropout)	(None, 112)	0
dense_102 (Dense)	(None, 54)	6102
batch_normalization_57 (Batch Normalization)	(None, 54)	216
dropout_37 (Dropout)	(None, 54)	0
dense_103 (Dense)	(None, 28)	1540
batch_normalization_58 (Batch Normalization)	(None, 28)	112
dropout_38 (Dropout)	(None, 28)	0
dense_104 (Dense)	(None, 16)	464
batch_normalization_59 (Batch Normalization)	(None, 16)	64
dropout_39 (Dropout)	(None, 16)	0

```
dense_105 (Dense)                (None, 10)                170
=====
Total params: 340,796
Trainable params: 339,640
Non-trainable params: 1,156
```

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 185us/step - loss: 1.7393 - acc: 0.4098 - val_loss: 0.6295 - val_acc: 0.8676

Epoch 2/20

60000/60000 [=====] - 6s 95us/step - loss: 0.9038 - acc: 0.7119 - val_loss: 0.2672 - val_acc: 0.9344

Epoch 3/20

60000/60000 [=====] - 6s 94us/step - loss: 0.5987 - acc: 0.8227 - val_loss: 0.1998 - val_acc: 0.9473

Epoch 4/20

60000/60000 [=====] - 6s 96us/step - loss: 0.4686 - acc: 0.8728 - val_loss: 0.1606 - val_acc: 0.9600

Epoch 5/20

60000/60000 [=====] - 6s 94us/step - loss: 0.4086 - acc: 0.8923 - val_loss: 0.1498 - val_acc: 0.9634

Epoch 6/20

60000/60000 [=====] - 6s 96us/step - loss: 0.3552 - acc: 0.9100 - val_loss: 0.1305 - val_acc: 0.9682

Epoch 7/20

60000/60000 [=====] - 6s 93us/step - loss: 0.3229 - acc: 0.9191 - val_loss: 0.1266 - val_acc: 0.9688

Epoch 8/20

60000/60000 [=====] - 6s 93us/step - loss: 0.2989 - acc: 0.9263 - val_loss: 0.1205 - val_acc: 0.9713

Epoch 9/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2848 - acc: 0.9300 - val_loss: 0.1136 - val_acc: 0.9721

Epoch 10/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2

```

591 - acc: 0.9367 - val_loss: 0.1105 - val_acc: 0.9748
Epoch 11/20
60000/60000 [=====] - 6s 94us/step - loss: 0.2
535 - acc: 0.9405 - val_loss: 0.1090 - val_acc: 0.9757
Epoch 12/20
60000/60000 [=====] - 6s 94us/step - loss: 0.2
362 - acc: 0.9430 - val_loss: 0.1039 - val_acc: 0.9757
Epoch 13/20
60000/60000 [=====] - 6s 95us/step - loss: 0.2
301 - acc: 0.9460 - val_loss: 0.0982 - val_acc: 0.9779
Epoch 14/20
60000/60000 [=====] - 6s 95us/step - loss: 0.2
122 - acc: 0.9485 - val_loss: 0.1020 - val_acc: 0.9780
Epoch 15/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2
077 - acc: 0.9505 - val_loss: 0.0991 - val_acc: 0.9790
Epoch 16/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2
069 - acc: 0.9517 - val_loss: 0.1021 - val_acc: 0.9781
Epoch 17/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1
980 - acc: 0.9531 - val_loss: 0.1009 - val_acc: 0.9800
Epoch 18/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1
881 - acc: 0.9557 - val_loss: 0.0942 - val_acc: 0.9792
Epoch 19/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1
875 - acc: 0.9569 - val_loss: 0.0998 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1
820 - acc: 0.9572 - val_loss: 0.0960 - val_acc: 0.9797

```

```

In [102]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
          print ('Test score:', score[0])
          print ('Test accuracy:', score[1])

          accuracy_3c= score[1]

          fig, ax = plt.subplots(1,1)

```

```


ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09595172328427434

Test accuracy: 0.9797



```

In [103]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Layer 1')


plt.subplot(1, 6, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='g')
plt.xlabel('Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 (d) MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 5 hidden layers

```
In [104]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.073)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
```

```

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean =0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(28, activation='relu', kernel_initializer=RandomNormal(mean =0.0, stddev=0.267, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean =0.0, stddev=0.353, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
print('\n')
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_106 (Dense)	(None, 368)	288880
batch_normalization_60 (Batch Normalization)	(None, 368)	1472
dropout_40 (Dropout)	(None, 368)	0
dense_107 (Dense)	(None, 112)	41328
batch_normalization_61 (Batch Normalization)	(None, 112)	448
dropout_41 (Dropout)	(None, 112)	0

dense_108 (Dense)	(None, 54)	6102
batch_normalization_62 (Batch Normalization)	(None, 54)	216
dropout_42 (Dropout)	(None, 54)	0
dense_109 (Dense)	(None, 28)	1540
batch_normalization_63 (Batch Normalization)	(None, 28)	112
dropout_43 (Dropout)	(None, 28)	0
dense_110 (Dense)	(None, 16)	464
batch_normalization_64 (Batch Normalization)	(None, 16)	64
dropout_44 (Dropout)	(None, 16)	0
dense_111 (Dense)	(None, 10)	170
=====		
Total params: 340,796		
Trainable params: 339,640		
Non-trainable params: 1,156		
None		

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 191us/step - loss: 2.1301 - acc: 0.2682 - val_loss: 1.0570 - val_acc: 0.6764
Epoch 2/20
60000/60000 [=====] - 6s 93us/step - loss: 1.3601 - acc: 0.5146 - val_loss: 0.6491 - val_acc: 0.8335
Epoch 3/20
60000/60000 [=====] - 6s 97us/step - loss: 1.0379 - acc: 0.6367 - val_loss: 0.4739 - val_acc: 0.9047
Epoch 4/20
60000/60000 [=====] - 6s 100us/step - loss: 0.

```

```
8556 - acc: 0.7115 - val_loss: 0.3282 - val_acc: 0.9279
Epoch 5/20
60000/60000 [=====] - 6s 98us/step - loss: 0.7
351 - acc: 0.7641 - val_loss: 0.2309 - val_acc: 0.9501
Epoch 6/20
60000/60000 [=====] - 6s 94us/step - loss: 0.6
363 - acc: 0.8065 - val_loss: 0.1850 - val_acc: 0.9536
Epoch 7/20
60000/60000 [=====] - 6s 94us/step - loss: 0.5
664 - acc: 0.8297 - val_loss: 0.1596 - val_acc: 0.9617
Epoch 8/20
60000/60000 [=====] - 6s 94us/step - loss: 0.5
100 - acc: 0.8508 - val_loss: 0.1647 - val_acc: 0.9613
Epoch 9/20
60000/60000 [=====] - 6s 93us/step - loss: 0.4
821 - acc: 0.8604 - val_loss: 0.1511 - val_acc: 0.9650
Epoch 10/20
60000/60000 [=====] - 6s 94us/step - loss: 0.4
623 - acc: 0.8682 - val_loss: 0.1405 - val_acc: 0.9682
Epoch 11/20
60000/60000 [=====] - 6s 93us/step - loss: 0.4
424 - acc: 0.8740 - val_loss: 0.1375 - val_acc: 0.9687
Epoch 12/20
60000/60000 [=====] - 6s 95us/step - loss: 0.4
197 - acc: 0.8821 - val_loss: 0.1242 - val_acc: 0.9713
Epoch 13/20
60000/60000 [=====] - 6s 95us/step - loss: 0.4
048 - acc: 0.8871 - val_loss: 0.1237 - val_acc: 0.9718
Epoch 14/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3
879 - acc: 0.8919 - val_loss: 0.1340 - val_acc: 0.9702
Epoch 15/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3
706 - acc: 0.8962 - val_loss: 0.1218 - val_acc: 0.9739
Epoch 16/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3
696 - acc: 0.8964 - val_loss: 0.1258 - val_acc: 0.9737
Epoch 17/20
60000/60000 [=====] - 6s 95us/step - loss: 0.3
```

```
624 - acc: 0.8993 - val_loss: 0.1280 - val_acc: 0.9734
Epoch 18/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3
470 - acc: 0.9017 - val_loss: 0.1154 - val_acc: 0.9758
Epoch 19/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3
434 - acc: 0.9042 - val_loss: 0.1138 - val_acc: 0.9767
Epoch 20/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3
339 - acc: 0.9066 - val_loss: 0.1138 - val_acc: 0.9774
```

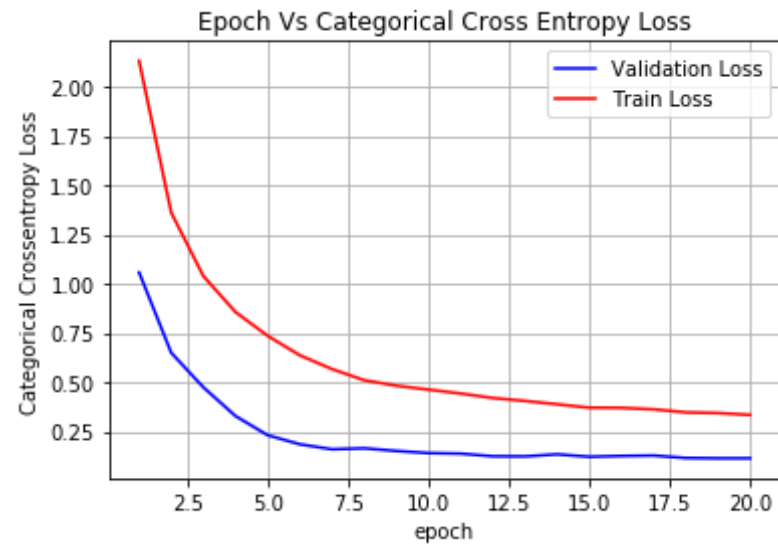
```
In [105]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print ('Test score:', score[0])
print ('Test accuracy:', score[1])
```

```
accuracy_3d= score[1]
```

```
fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')
```

```
#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.11384864466506987
Test accuracy: 0.9774
```



```
In [106]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Layer 1')

plt.subplot(1, 6, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 2 ')
```

```


plt.subplot(1, 6, 3)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='g')
plt.xlabel('Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



3 (e) MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 5 hidden layers

```
In [107]: model_relu = Sequential()

# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0,\sigma) = N(0,0.073)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0,\sigma) = N(0,0.133)$ 

model_relu.add(Dense(368, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.133, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(54, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.192, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(28, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.267, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.353, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.6))

model_relu.add(Dense(output_dim, activation='softmax'))
```

```

print(model_relu.summary())
print('\n')
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_112 (Dense)	(None, 368)	288880
batch_normalization_65 (Batch Normalization)	(None, 368)	1472
dropout_45 (Dropout)	(None, 368)	0
dense_113 (Dense)	(None, 112)	41328
batch_normalization_66 (Batch Normalization)	(None, 112)	448
dropout_46 (Dropout)	(None, 112)	0
dense_114 (Dense)	(None, 54)	6102
batch_normalization_67 (Batch Normalization)	(None, 54)	216
dropout_47 (Dropout)	(None, 54)	0
dense_115 (Dense)	(None, 28)	1540
batch_normalization_68 (Batch Normalization)	(None, 28)	112
dropout_48 (Dropout)	(None, 28)	0
dense_116 (Dense)	(None, 16)	464
batch_normalization_69 (Batch Normalization)	(None, 16)	64
dropout_49 (Dropout)	(None, 16)	0

```
dense_117 (Dense)                (None, 10)                170
=====
Total params: 340,796
Trainable params: 339,640
Non-trainable params: 1,156
```

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 204us/step - loss: 2.4759 - acc: 0.1602 - val_loss: 1.7847 - val_acc: 0.4354

Epoch 2/20

60000/60000 [=====] - 6s 95us/step - loss: 1.9158 - acc: 0.2868 - val_loss: 1.4256 - val_acc: 0.6718

Epoch 3/20

60000/60000 [=====] - 6s 95us/step - loss: 1.6487 - acc: 0.3883 - val_loss: 1.0411 - val_acc: 0.7306

Epoch 4/20

60000/60000 [=====] - 6s 97us/step - loss: 1.4051 - acc: 0.4742 - val_loss: 0.7940 - val_acc: 0.7484

Epoch 5/20

60000/60000 [=====] - 6s 97us/step - loss: 1.2284 - acc: 0.5322 - val_loss: 0.6732 - val_acc: 0.8191

Epoch 6/20

60000/60000 [=====] - 6s 94us/step - loss: 1.1199 - acc: 0.5722 - val_loss: 0.6110 - val_acc: 0.8291

Epoch 7/20

60000/60000 [=====] - 6s 94us/step - loss: 1.0372 - acc: 0.6057 - val_loss: 0.5386 - val_acc: 0.8503

Epoch 8/20

60000/60000 [=====] - 6s 95us/step - loss: 0.9872 - acc: 0.6276 - val_loss: 0.5028 - val_acc: 0.8442

Epoch 9/20

60000/60000 [=====] - 6s 94us/step - loss: 0.9467 - acc: 0.6433 - val_loss: 0.4719 - val_acc: 0.8572

Epoch 10/20

60000/60000 [=====] - 6s 94us/step - loss: 0.9


```

038 - acc: 0.6594 - val_loss: 0.4472 - val_acc: 0.8301
Epoch 11/20
60000/60000 [=====] - 6s 95us/step - loss: 0.8
713 - acc: 0.6673 - val_loss: 0.4292 - val_acc: 0.8188
Epoch 12/20
60000/60000 [=====] - 6s 94us/step - loss: 0.8
404 - acc: 0.6760 - val_loss: 0.4225 - val_acc: 0.8353
Epoch 13/20
60000/60000 [=====] - 6s 95us/step - loss: 0.8
349 - acc: 0.6808 - val_loss: 0.4083 - val_acc: 0.8536
Epoch 14/20
60000/60000 [=====] - 6s 96us/step - loss: 0.7
990 - acc: 0.6894 - val_loss: 0.4089 - val_acc: 0.8408
Epoch 15/20
60000/60000 [=====] - 6s 96us/step - loss: 0.7
921 - acc: 0.6915 - val_loss: 0.3960 - val_acc: 0.8587
Epoch 16/20
60000/60000 [=====] - 6s 94us/step - loss: 0.7
742 - acc: 0.7013 - val_loss: 0.3867 - val_acc: 0.9119
Epoch 17/20
60000/60000 [=====] - 6s 96us/step - loss: 0.7
657 - acc: 0.7036 - val_loss: 0.3724 - val_acc: 0.8687
Epoch 18/20
60000/60000 [=====] - 6s 95us/step - loss: 0.7
468 - acc: 0.7080 - val_loss: 0.3733 - val_acc: 0.9071
Epoch 19/20
60000/60000 [=====] - 6s 94us/step - loss: 0.7
454 - acc: 0.7083 - val_loss: 0.3576 - val_acc: 0.8720
Epoch 20/20
60000/60000 [=====] - 6s 95us/step - loss: 0.7
385 - acc: 0.7146 - val_loss: 0.3569 - val_acc: 0.8666

```

```

In [108]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
          print ('Test score:', score[0])
          print ('Test accuracy:', score[1])

          accuracy_3e= score[1]

          fig, ax = plt.subplots(1,1)

```

```


ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

#List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.35688953695297243

Test accuracy: 0.8666



```

In [109]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Layer 1')


plt.subplot(1, 6, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='g')
plt.xlabel('Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



4 a) Procedure Followed:

STEP 1: Load the MNIST dataset, shuffle the data and split it between train and test.

STEP 2: Normalize the input data and one-hot encode class label data before feeding into the model. This conversion is needed for Multi layer perceptron.

STEP 3: Add the hidden layers and provide the required input parameters like number of neurons for each layer, activation unit, input shape and weight initialization with kernel_initializer.

STEP 4: Compile the model and fit it on training data.

STEP 5: Evaluate the model on test data and plot the error plots for each epoch.

STEP 6: Get the weights from trained model and plot them to see their distribution.

STEP 7: Repeat the steps from 3 to 6 for hidden layers 2, 3, 5 and try the cases with batch normalization and different dropout for each of the architectures.

STEP 8: Summarize the test accuracy for each model.

4 b) Models performance comparison with different architectures

```
In [110]: # comparing Models performance using Prettytable library
from prettytable import PrettyTable
x = PrettyTable()

# Name of the models
models = ['MLP + ReLU + ADAM with 2 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization with 2 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 2 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 2 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 2 hidden layers',
          'MLP + ReLU + ADAM with 3 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization with 3 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 3 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 3 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 3 hidden layers',
          'MLP + ReLU + ADAM with 5 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization with 5 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 5 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 5 hidden layers',
          'MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 5 hidden layers'
          ]
```

```

# Test Accuracy
test_accuracy = [accuracy_1a,accuracy_1b,accuracy_1c,accuracy_1d,accuracy_1e,accuracy_2a,accuracy_2b,accuracy_2c,accuracy_2d,accuracy_2e,accuracy_3a,accuracy_3b,accuracy_3c,accuracy_3d,accuracy_3e]
sno = [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]

# Adding columns
x.add_column("S.NO.",sno)
x.add_column("MLP MODEL",models)
x.add_column("Test Accuracy",test_accuracy)

# Printing the Table
print(x)

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| S.NO. |                                     MLP MODEL                                     |
|       | | Test Accuracy | |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|  1  | | MLP + ReLU + ADAM with 2 hidden layers | | 0.9815 | |
|  2  | | MLP + ReLU + ADAM + Batch Normalization with 2 hidden layers | | 0.9834 | |
|  3  | | MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 2 hidden layers | | 0.9827 | |
|  4  | | MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 2 hidden layers | | 0.9834 | |
|  5  | | MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 2 hidden layers | | 0.9811 | |
|  6  | | MLP + ReLU + ADAM with 3 hidden layers | | 0.9791 | |
|  7  | | MLP + ReLU + ADAM + Batch Normalization with 3 hidden layers | | 0.979 | |
|  8  | | MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 3 hidden layers | | 0.9821 | |
|  9  | | MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 3 hidden layers | | 0.9804 | |
| 10  | | MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 3 hidden layers | | 0.9786 | |

```

11	MLP + ReLU + ADAM with 5 hidden layers	
	0.9763	
12	MLP + ReLU + ADAM + Batch Normalization with 5 hidden layers	
	0.9794	
13	MLP + ReLU + ADAM + Batch Normalization + Dropout(0.4) with 5 hidden layers	
	0.9797	
14	MLP + ReLU + ADAM + Batch Normalization + Dropout(0.5) with 5 hidden layers	
	0.9774	
15	MLP + ReLU + ADAM + Batch Normalization + Dropout(0.6) with 5 hidden layers	
	0.8666	
+-----+		
-----+		

Observation:

- 1) It is observed that accuracy decreased slightly when dropout is more than 0.5.
- 2) Accuracy increased when batch normalization and dropout is added to the models.